

Project ML - Part II

Group member name and IDs

Asifa Batool ----- F2023266753

Mariha Javed ----- F2023266780

Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score, precision_score
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, roc_auc_score, roc_curve
import warnings
warnings.filterwarnings('ignore')

# Load dataset
df = pd.read_csv("/content/creditcard_fraud_detection.csv")
df.head()

{"type": "dataframe", "variable_name": "df"}
```

Checking Shape, Types, Nulls

```
print("Shape of dataset:", df.shape)
print("\nData types:\n", df.dtypes)
print("\nMissing values:\n", df.isnull().sum())
```

Shape of dataset: (49610, 31)

Data types:

Time	int64
V1	float64
V2	float64
V3	float64
V4	float64
V5	float64
V6	float64

```
V7      float64
V8      float64
V9      float64
V10     float64
V11     float64
V12     float64
V13     float64
V14     float64
V15     float64
V16     float64
V17     float64
V18     float64
V19     float64
V20     float64
V21     float64
V22     float64
V23     float64
V24     float64
V25     float64
V26     float64
V27     float64
V28     float64
Amount  float64
Class   float64
dtype: object
```

Missing values:

```
Time    0
V1      0
V2      0
V3      0
V4      1
V5      1
V6      1
V7      1
V8      1
V9      1
V10     1
V11     1
V12     1
V13     1
V14     1
V15     1
V16     1
V17     1
V18     1
V19     1
V20     1
V21     1
```

```
V22      1
V23      1
V24      1
V25      1
V26      1
V27      1
V28      1
Amount   1
Class    1
dtype: int64
```

```
# removing null values
df.dropna(inplace=True)
```

Dataset Summary and plotting target feature

```
print(df.describe())
print("\nClass Distribution:\n", df['Class'].value_counts())
sns.countplot(data=df, x='Class')
plt.title('Class Distribution (0 = Not Fraud, 1 = Fraud)')
plt.show()
```

	Time	V1	V2	V3
V4 \				
count	49609.000000	49609.000000	49609.000000	49609.000000
mean	28803.247193	-0.242479	0.012321	0.692985
std	13097.419648	1.885778	1.630608	1.510566
min	0.000000	-56.407510	-72.715728	-32.965346
25%	21734.000000	-0.992814	-0.562958	0.217595
50%	33390.000000	-0.247067	0.079334	0.796997
75%	38852.000000	1.155641	0.732318	1.430964
max	44134.000000	1.960497	18.183626	4.101716

	V5	V6	V7	V8
V9 \				
count	49609.000000	49609.000000	49609.000000	49609.000000
mean	-0.257016	0.104114	-0.120255	0.053442
std	1.413057	1.310705	1.283507	1.224245
min	-42.147898	-26.160506	-26.548144	-41.484823

9.283925					
25%	-0.866471	-0.635669	-0.605928	-0.146749	-
0.611499					
50%	-0.287810	-0.150940	-0.076595	0.058406	
0.012150					
75%	0.283513	0.493918	0.424969	0.331555	
0.819242					
max	34.801666	22.529298	36.677268	20.007208	
10.392889					

		V21	V22	V23	V24 \
count	...	49609.000000	49609.000000	49609.000000	49609.000000
mean	...	-0.028396	-0.107154	-0.040123	0.007997
std	...	0.736050	0.637733	0.590810	0.594121
min	...	-20.262054	-8.593642	-26.751119	-2.836627
25%	...	-0.231664	-0.529531	-0.179110	-0.322243
50%	...	-0.068396	-0.082137	-0.051560	0.061999
75%	...	0.108082	0.307262	0.078474	0.401392
max	...	22.614889	5.805795	17.297845	4.014444

	V25	V26	V27	V28
Amount \				
count	49609.000000	49609.000000	49609.000000	49609.000000
49609.000000				
mean	0.135954	0.020813	0.004792	0.004533
93.120688				
std	0.439067	0.501438	0.388364	0.333225
253.265971				
min	-7.495741	-1.577118	-8.567638	-9.617915
0.000000				
25%	-0.127983	-0.330532	-0.063339	-0.006675
7.610000				
50%	0.175766	-0.071826	0.008986	0.022155
25.000000				
75%	0.421960	0.300180	0.083910	0.076342
85.000000				
max	5.525093	3.517346	11.135740	33.847808
12910.930000				


	Class
count	49609.000000
mean	0.002983
std	0.054539
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 31 columns]

```
Class Distribution:
Class
0.0    49461
1.0     148
Name: count, dtype: int64
```

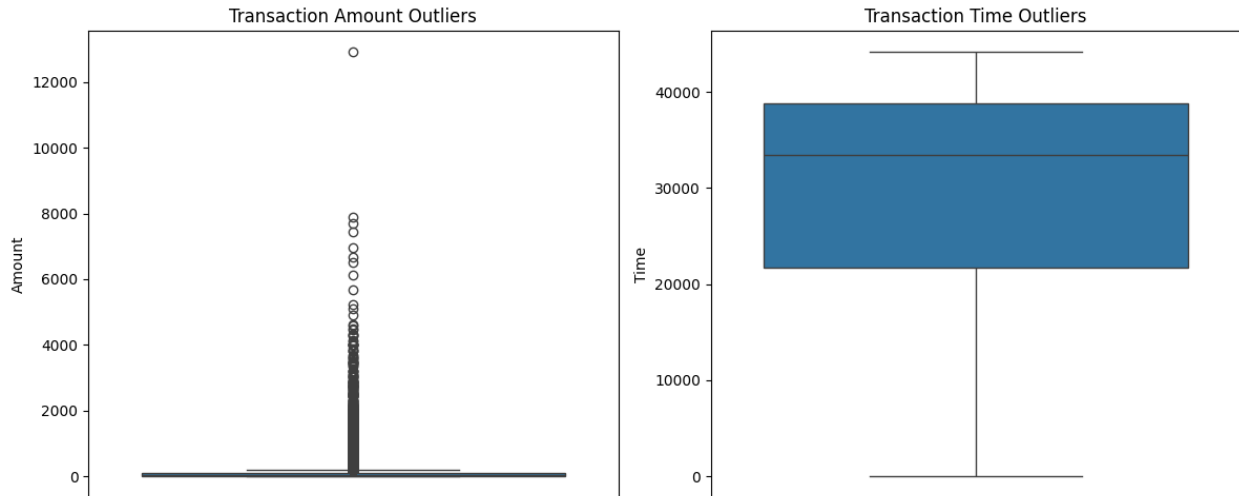


```
# Heatmap
plt.figure(figsize=(16, 10))
sns.heatmap(df.corr(), cmap='coolwarm', linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()
```



Check for Outliers in 'Amount' and 'Time'

```
fig, axs = plt.subplots(1, 2, figsize=(12, 5))
sns.boxplot(data=df, y='Amount', ax=axs[0])
axs[0].set_title("Transaction Amount Outliers")
sns.boxplot(data=df, y='Time', ax=axs[1])
axs[1].set_title("Transaction Time Outliers")
plt.tight_layout()
plt.show()
```



Feature Scaling - StandardScaler for 'Time' and 'Amount'

```
scaler = StandardScaler()
df['scaled_amount'] = scaler.fit_transform(df[['Amount']])
df['scaled_time'] = scaler.fit_transform(df[['Time']])
df.drop(['Time', 'Amount'], axis=1, inplace=True)
df.head()
```

```
{"type": "dataframe", "variable_name": "df"}
```

Divide dataset using Train-Test Split

```
# Features and target
X = df.drop('Class', axis=1)
y = df['Class']

# Stratified split to maintain class ratio
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, stratify=y, random_state=42)

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)
```

```
Training set shape: (39687, 30)
Testing set shape: (9922, 30)
```

Import ML models

```
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
```

```
# function for evaluate the model
def evaluate_model(y_true, y_pred, model_name):
    print(f"----- {model_name} Evaluation -----")
    print("Accuracy:", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred))
    print("Recall:", recall_score(y_true, y_pred))
    print("F1 Score:", f1_score(y_true, y_pred))
    print("\nConfusion Matrix:\n", confusion_matrix(y_true, y_pred))
    print("\nClassification Report:\n", classification_report(y_true,
y_pred))
    print("-" * 40)
```

LogisticRegression

```
lr = LogisticRegression()
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
evaluate_model(y_test, y_pred_lr, "Logistic Regression")
```

----- Logistic Regression Evaluation -----

Accuracy: 0.9979842773634348

Precision: 0.65625

Recall: 0.7

F1 Score: 0.6774193548387096

Confusion Matrix:

```
[[9881  11]
 [   9  21]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	9892
1.0	0.66	0.70	0.68	30
accuracy			1.00	9922
macro avg	0.83	0.85	0.84	9922
weighted avg	1.00	1.00	1.00	9922

Decision Tree

```
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
evaluate_model(y_test, y_pred_dt, "Decision Tree")
```

----- Decision Tree Evaluation -----

Accuracy: 0.9990929248135456


```
Precision: 0.92
Recall: 0.7666666666666667
F1 Score: 0.8363636363636363
```

```
Confusion Matrix:
```

```
[[9890   2]
 [   7  23]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	9892
1.0	0.92	0.77	0.84	30
accuracy			1.00	9922
macro avg	0.96	0.88	0.92	9922
weighted avg	1.00	1.00	1.00	9922

```
-----
```

Random Forest

```
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
evaluate_model(y_test, y_pred_rf, "Random Forest")
```

```
----- Random Forest Evaluation -----
```

```
Accuracy: 0.9992944970772022
Precision: 0.9259259259259259
Recall: 0.8333333333333334
F1 Score: 0.8771929824561403
```

```
Confusion Matrix:
```

```
[[9890   2]
 [   5  25]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	9892
1.0	0.93	0.83	0.88	30
accuracy			1.00	9922
macro avg	0.96	0.92	0.94	9922
weighted avg	1.00	1.00	1.00	9922

```
-----
```

KNN

```
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)
evaluate_model(y_test, y_pred_knn, "K-Nearest Neighbors")
```

----- K-Nearest Neighbors Evaluation -----

Accuracy: 0.9993952832090305
Precision: 0.9285714285714286
Recall: 0.8666666666666667
F1 Score: 0.896551724137931

Confusion Matrix:

```
[[9890   2]
 [   4  26]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	9892
1.0	0.93	0.87	0.90	30
accuracy			1.00	9922
macro avg	0.96	0.93	0.95	9922
weighted avg	1.00	1.00	1.00	9922

Support Vector Machine (SVM)

```
svm = SVC(kernel='rbf', random_state=42)
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)
evaluate_model(y_test, y_pred_svm, "Support Vector Machine")
```

----- Support Vector Machine Evaluation -----

Accuracy: 0.9993952832090305
Precision: 0.9615384615384616
Recall: 0.8333333333333334
F1 Score: 0.8928571428571429

Confusion Matrix:

```
[[9891   1]
 [   5  25]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	9892

	1.0	0.96	0.83	0.89	30
accuracy				1.00	9922
macro avg	0.98	0.92	0.95		9922
weighted avg	1.00	1.00	1.00		9922

Naive Bayes

```
nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred_nb = nb.predict(X_test)
evaluate_model(y_test, y_pred_nb, "Naive Bayes")
```

----- Naive Bayes Evaluation -----

Accuracy: 0.9822616407982262
Precision: 0.13131313131313133
Recall: 0.8666666666666667
F1 Score: 0.22807017543859648

Confusion Matrix:

```
[[9720  172]
 [    4   26]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	1.00	0.98	0.99	9892
1.0	0.13	0.87	0.23	30
accuracy			0.98	9922
macro avg	0.57	0.92	0.61	9922
weighted avg	1.00	0.98	0.99	9922

ANN

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset, random_split

# Convert to tensors
X_train_tensor = torch.tensor(X_train.values, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train.values,
dtype=torch.float32).view(-1, 1)

X_test_tensor = torch.tensor(X_test.values, dtype=torch.float32)
```

```

y_test_tensor = torch.tensor(y_test.values,
dtype=torch.float32).view(-1, 1)

# Create datasets and data loaders
train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
test_dataset = TensorDataset(X_test_tensor, y_test_tensor)

train_loader = DataLoader(train_dataset, batch_size=2048,
shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=2048, shuffle=False)

class ANNModel(nn.Module):
    def __init__(self, input_dim):
        super(ANNModel, self).__init__()
        self.fc1 = nn.Linear(input_dim, 32)
        self.dropout1 = nn.Dropout(0.2)
        self.fc2 = nn.Linear(32, 16)
        self.dropout2 = nn.Dropout(0.2)
        self.output = nn.Linear(16, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = self.dropout1(x)
        x = torch.relu(self.fc2(x))
        x = self.dropout2(x)
        x = self.sigmoid(self.output(x))
        return x

input_dim = X_train.shape[1]
model = ANNModel(input_dim)

criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

num_epochs = 20
for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    for inputs, labels in train_loader:
        outputs = model(inputs)
        loss = criterion(outputs, labels)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}/{num_epochs}, Loss:
{running_loss/len(train_loader):.4f}")

```

```
Epoch 1/20, Loss: 0.5459
Epoch 2/20, Loss: 0.4349
Epoch 3/20, Loss: 0.2950
Epoch 4/20, Loss: 0.1604
Epoch 5/20, Loss: 0.0801
Epoch 6/20, Loss: 0.0443
Epoch 7/20, Loss: 0.0284
Epoch 8/20, Loss: 0.0210
Epoch 9/20, Loss: 0.0164
Epoch 10/20, Loss: 0.0142
Epoch 11/20, Loss: 0.0118
Epoch 12/20, Loss: 0.0108
Epoch 13/20, Loss: 0.0097
Epoch 14/20, Loss: 0.0090
Epoch 15/20, Loss: 0.0086
Epoch 16/20, Loss: 0.0080
Epoch 17/20, Loss: 0.0079
Epoch 18/20, Loss: 0.0077
Epoch 19/20, Loss: 0.0073
Epoch 20/20, Loss: 0.0070
```

```
model.eval()
with torch.no_grad():
    y_pred_proba = model(X_test_tensor).numpy()
    y_pred = (y_pred_proba > 0.5).astype(int)

evaluate_model(y_test, y_pred, "PyTorch ANN")
```

----- PyTorch ANN Evaluation -----

Accuracy: 0.9969764160451522
Precision: 0.0
Recall: 0.0
F1 Score: 0.0

Confusion Matrix:

```
[[9892   0]
 [  30   0]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	9892
1.0	0.00	0.00	0.00	30
accuracy			1.00	9922
macro avg	0.50	0.50	0.50	9922
weighted avg	0.99	1.00	1.00	9922

```

fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
auc_score = roc_auc_score(y_test, y_pred_proba)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f"PyTorch ANN (AUC = {auc_score:.4f})",
color='green')
plt.plot([0, 1], [0, 1], 'k--')
plt.title("PyTorch ANN ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid(True)
plt.show()

```

